

Laporan Eksperimen LoRA

Juyus Muhammad Adinulhaq
25/572923/PPA/07202
Kapita Selektta Ilmu Komputer

Repository GitHub: <https://github.com/haqiqien/LoRA>

Notebook yang digunakan dalam laporan ini adalah `section4.ipynb`, `LoRA-default.ipynb`, dan `LoRA-exp.ipynb`.

1 Pendahuluan

LoRA (Low-Rank Adaptation) adalah cara melakukan fine-tuning tanpa mengubah semua bobot model. Bobot model dasar dibekukan, kemudian ditambahkan matriks kecil yang dilatih pada layer tertentu. Sebaliknya, full fine-tuning mengubah seluruh bobot model sehingga biasanya membutuhkan lebih banyak memori, waktu, dan ruang penyimpanan. Dengan LoRA, adapter dapat disimpan secara terpisah dan digabungkan dengan model dasar jika diperlukan. Cara ini lebih ringan, tetapi hasilnya tetap dipengaruhi oleh rank, alpha, dropout, layer target, kualitas data, dan jumlah epoch. Pada tugas ini, LoRA digunakan untuk menyesuaikan HuggingFaceTB/SmolLM2-135M dengan dataset percakapan HuggingFaceTB/smoltalk.

2 Peran dan Perubahan Notebook

Ketiga notebook memiliki fungsi yang berbeda:

File	Peran
<code>section4.ipynb</code>	Notebook original dari Google Classroom. File ini menjadi dasar pengerjaan, tetapi mengalami beberapa error ketika dijalankan di Google Colab.
<code>LoRA-default.ipynb</code>	Notebook fine-tuning utama versi perbaikan. File ini dibuat agar instalasi, training, penyimpanan, merge, dan inference dapat berjalan.
<code>LoRA-exp.ipynb</code>	Notebook eksperimen hyperparameter LoRA. File ini menjalankan beberapa run dengan parameter berbeda lalu membandingkan training loss.

Beberapa perbaikan utama dari notebook original adalah sebagai berikut.

Bagian	Masalah awal	Perbaikan
Package	Versi <code>torchao</code> tidak cocok dengan <code>peft</code> .	Memasang <code>torchao>=0.16.0</code> .
Warmup	Langsung memakai <code>warmup_ratio</code> .	Memilih <code>warmup_ratio</code> atau <code>warmup_steps</code> sesuai versi TRL.
Precision	<code>bf16=True</code> selalu aktif.	BF16 hanya aktif jika hardware mendukung.
Trainer API	Memakai nama argumen lama secara langsung.	Memeriksa dukungan <code>max_length</code> , <code>max_seq_length</code> , <code>tokenizer</code> , dan <code>processing_class</code> .
Chat template	Menganggap tokenizer sudah memiliki template.	Menambahkan fallback chat template.
Monitoring	Hanya memanggil <code>trainer.train()</code> .	Menampilkan <code>step</code> , <code>loss</code> , <code>epoch</code> , <code>learning rate</code> , waktu, dan ETA.
Model path	Folder lokal dianggap sebagai repository Hugging Face.	Menggunakan folder adapter dan hasil merge lokal yang benar.

Perubahan diatas terutama diperlukan agar notebook bisa mengikuti versi library yang tersedia di Google Colab.

3 Eksperimen

Konfigurasi umum yang digunakan adalah satu epoch, learning rate $2e-4$, batch size 2, gradient accumulation 2, dan sequence length 1512. Notebook default memakai rank 6, alpha 8, dropout 0.05, dan target `all-linear`. Pada notebook eksperimen, parameter tersebut diubah pada setiap run.

Run	Rank	Alpha	Dropout	Target	Loss
Default	6	8	0.05	<code>all-linear</code>	1.8389
<code>r4_a8_d005_all</code>	4	8	0.05	<code>all-linear</code>	1.8291
<code>r8_a16_d005_all</code>	8	16	0.05	<code>all-linear</code>	1.6855
<code>r8_a16_d010_qv</code>	8	16	0.10	<code>q_proj,v_proj</code>	2.1138
<code>r16_a32_d010_qv</code>	16	32	0.10	<code>q_proj,v_proj</code>	1.9363

Tabel 3: Parameter dan training loss setiap run.

Run dengan hasil loss paling rendah adalah `r8_a16_d005_all`. Namun, nilai ini hanya berasal dari data training karena eksperimen belum memakai validation split.

4 Hasil Inference

Prompt	Sebelum <i>fine-tuning</i>	Sesudah <i>fine-tuning</i>
Ibu kota Jerman	Belum tercatat pada notebook awal.	Model menjawab Berlin, tetapi kemudian mencampur percakapan dan fakta lain.
Fungsi factorial Python	Belum tercatat.	Model menjelaskan factorial, tetapi tidak memberikan kode Python.
Pagar 25 x 15 kaki	Belum tercatat.	Model menjawab 25 kaki, padahal jawaban yang benar adalah $2(25 + 15) = 80$ kaki.
Buah dan sayur	Belum tercatat.	Jawaban masih berhubungan dengan prompt, tetapi definisinya kabur dan berulang.

5 Analisis

Rank 8 dengan alpha 16 dan target `all-linear` memberi training loss terendah. Target `q_proj` dan `v_proj` berjalan lebih cepat karena layer yang dilatih lebih sedikit, tetapi loss-nya lebih tinggi. Ini menunjukkan trade-off antara efisiensi dan kapasitas adaptasi.

Walaupun training loss turun, kualitas jawaban belum stabil. Model masih mengulang teks, mencampur percakapan, gagal mengikuti instruksi untuk membuat kode factorial, dan salah menghitung keliling. Jadi, training loss belum cukup untuk menilai kemampuan instruction-following. Kemungkinan penyebabnya adalah ukuran model yang kecil, training hanya satu epoch, kualitas dataset, format chat template, dan pengaturan generation. Eksperimen lanjutan sebaiknya memakai validation split, mencatat `eval_loss`, menyimpan output baseline, mengontrol temperature, dan menggunakan `max_new_tokens` secara eksplisit.

6 Dokumentasi Screenshot

Screenshot disimpan pada folder `screenshots/`.

```
[3]
15s
# Import necessary libraries
from transformers import AutoModelForCausalLM, AutoTokenizer
from datasets import load_dataset
from trl import SFTConfig, SFTTrainer, setup_chat_format
import torch

device = (
    "cuda"
    if torch.cuda.is_available()
    else "mps" if torch.backends.mps.is_available() else "cpu"
)

# Load the model and tokenizer
model_name = "HuggingFaceTB/SmolLM2-135M"

model = AutoModelForCausalLM.from_pretrained(
    pretrained_model_name_or_path=model_name
).to(device)
tokenizer = AutoTokenizer.from_pretrained(pretrained_model_name_or_path=model_name)

# Set up the chat format
model, tokenizer = setup_chat_format(model=model, tokenizer=tokenizer)

# Set our name for the finetune to be saved &/ uploaded to
finetune_name = "SmolLM2-FT-MyDataset"
finetune_tags = ["smol-course", "module_1"]

-----
ModuleNotFoundError                                Traceback (most recent call last)
/tmp/ipykernel_2933/109771837.py in <cell line: 0>()
      2 from transformers import AutoModelForCausalLM, AutoTokenizer
      3 from datasets import load_dataset
----> 4 from trl import SFTConfig, SFTTrainer, setup_chat_format
      5 import torch
      6

ModuleNotFoundError: No module named 'trl'

-----
NOTE: If your import is failing due to a missing package, you can
manually install dependencies using either !pip or !apt.

To view examples of installing some common dependencies, click the
"Open Examples" button below.
-----
```

Gambar 1: Error pada notebook original.

```

# Training configuration
# Hyperparameters based on QLoRA paper recommendations
# TRL versions differ: older versions accept warmup_ratio, while newer
# versions expose warmup_steps instead.
import inspect
import math

if "warmup_ratio" in inspect.signature(SFTConfig).parameters:
    warmup_kwargs = {"warmup_ratio": 0.03}
else:
    updates_per_epoch = math.ceil(len(dataset["train"]) / (2 * 2))
    warmup_kwargs = {"warmup_steps": max(1, round(0.03 * updates_per_epoch))}

args = SFTConfig(
    # Output settings
    output_dir=finetune_name, # Directory to save model checkpoints
    # Training duration
    num_train_epochs=1, # Number of training epochs
    # Batch size settings
    per_device_train_batch_size=2, # Batch size per GPU
    gradient_accumulation_steps=2, # Accumulate gradients for larger effective batch
    # Memory optimization
    gradient_checkpointing=True, # Trade compute for memory savings
    # Optimizer settings
    optim="adamw_torch_fused", # Use fused AdamW for efficiency
    learning_rate=2e-4, # Learning rate (QLoRA paper)
    max_grad_norm=0.3, # Gradient clipping threshold
    # Learning rate schedule
    **warmup_kwargs, # 3% warmup, compatible with the installed TRL version
    lr_scheduler_type="constant", # Keep learning rate constant after warmup
    # Logging and saving
    logging_steps=1, # Log metrics every training step
    disable_tqdm=False, # Show the training progress bar
    save_strategy="epoch", # Save checkpoint every epoch
    # Precision settings
    # Enable bf16 only on supported CUDA hardware; use fp32 otherwise.
    bf16=torch.cuda.is_available() and torch.cuda.is_bf16_supported(),
    # Integration settings
    push_to_hub=False, # Don't push to HuggingFace Hub
    report_to="none", # Disable external logging
)

```

Gambar 2: Konfigurasi setelah perbaikan.

```
[ ] # Start training and display progress in the notebook.
from transformers import TrainerCallback
import time

class ConsoleProgressCallback(TrainerCallback):
    def on_train_begin(self, args, state, control, **kwargs):
        self.start_time = time.time()

    def on_log(self, args, state, control, logs=None, **kwargs):
        if logs:
            metrics = {key: value for key, value in logs.items() if key in {"loss", "learning_rate", "epoch"}}
            if metrics:
                self.last_metrics = metrics

    def on_step_end(self, args, state, control, **kwargs):
        if state.max_steps and state.global_step > 0:
            elapsed = time.time() - self.start_time
            steps_per_second = state.global_step / max(elapsed, 1e-6)
            remaining = (state.max_steps - state.global_step) / max(steps_per_second, 1e-6)
            percent = 100 * state.global_step / state.max_steps
            metrics = getattr(self, "last_metrics", {})
            loss = metrics.get("loss", "-")
            learning_rate = metrics.get("learning_rate", "-")
            print(
                f"Progress: {percent:6.2f}% | step {state.global_step}/{state.max_steps} | "
                f"epoch {state.epoch:2f} | loss {loss} | lr {learning_rate} | "
                f"elapsed {elapsed / 60:.1f} min | ETA {remaining / 60:.1f} min",
                flush=True,
            )

trainer.add_callback(ConsoleProgressCallback())
print(f"Starting training: {args.num_train_epochs} epoch(s), {len(dataset['train'])} training examples", flush=True)
train_result = trainer.train()

print("Training finished.")
print("Training metrics:", train_result.metrics)

# Save the trained LoRA adapter/model.
print(f"Saving model to: {args.output_dir}")
trainer.save_model()
print("Model saved successfully.")

# Display the recorded logs for monitoring after training.
trainer.state.log_history
```

```

Progress: 1.33% | step 1/75 | epoch 0.01 | loss - | lr - | elapsed 0.1 min | ETA 9.5 min
Progress: 2.67% | step 2/75 | epoch 0.03 | loss 2.740184783935547 | lr 0.0002 | elapsed 0.2 min | ETA 8.5 min
Progress: 4.00% | step 3/75 | epoch 0.04 | loss 2.595839023590088 | lr 0.0002 | elapsed 0.3 min | ETA 8.4 min
Progress: 5.33% | step 4/75 | epoch 0.05 | loss 2.6326775550842285 | lr 0.0002 | elapsed 0.5 min | ETA 8.1 min
Progress: 6.67% | step 5/75 | epoch 0.07 | loss 2.66498088366699 | lr 0.0002 | elapsed 0.6 min | ETA 8.0 min
Progress: 8.00% | step 6/75 | epoch 0.08 | loss 2.58721923828125 | lr 0.0002 | elapsed 0.7 min | ETA 7.9 min
Progress: 9.33% | step 7/75 | epoch 0.09 | loss 2.581766128540039 | lr 0.0002 | elapsed 0.8 min | ETA 7.8 min
Progress: 10.67% | step 8/75 | epoch 0.11 | loss 2.595880985260098 | lr 0.0002 | elapsed 0.9 min | ETA 7.6 min
Progress: 12.00% | step 9/75 | epoch 0.12 | loss 2.7096400260925293 | lr 0.0002 | elapsed 1.0 min | ETA 7.3 min
Progress: 13.33% | step 10/75 | epoch 0.13 | loss 2.46652889251709 | lr 0.0002 | elapsed 1.1 min | ETA 7.2 min
Progress: 14.67% | step 11/75 | epoch 0.15 | loss 2.4829750061035156 | lr 0.0002 | elapsed 1.2 min | ETA 7.2 min
Progress: 16.00% | step 12/75 | epoch 0.16 | loss 2.4867944717407227 | lr 0.0002 | elapsed 1.3 min | ETA 7.1 min

```

Gambar 3: Progress training.

	name	r	alpha	dropout	target_modules	output_dir	train_runtime	train_samples_per_second	train_steps_per_second	total_flos	train_loss	epoch	wall_time_minutes
1	r8_a16_d005_all	8	16	0.05	all-linear	SmolLM2-FT-MyDataset-exp-r8_a16_d005_all	513.0160	0.581	0.146	2.863064e+14	1.685517	1.0	8.56
0	r4_a8_d005_all	4	8	0.05	all-linear	SmolLM2-FT-MyDataset-exp-r4_a8_d005_all	512.2823	0.582	0.146	2.830885e+14	1.829089	1.0	8.56
3	r16_a32_d010_qv	16	32	0.10	[q_proj, v_proj]	SmolLM2-FT-MyDataset-exp-r16_a32_d010_qv	452.4809	0.659	0.166	2.822992e+14	1.936278	1.0	7.55
2	r8_a16_d010_qv	8	16	0.10	[q_proj, v_proj]	SmolLM2-FT-MyDataset-exp-r8_a16_d010_qv	452.1298	0.659	0.166	2.810848e+14	2.113827	1.0	7.55

Gambar 4: Tabel hasil eksperimen.

[illegible]

Gambar 5: Output inference, bagian pertama.

```

prompt:
A rectangular garden has a length of 25 feet and a width of 15 feet. If you want to build a fence around the entire garden, how many feet of fencing will you need?
[transformers] Both "max_new_tokens" (~256) and "max_length" (~20) seem to have been set. "max_new_tokens" will take precedence. Please refer to the documentation for more information. (https://huggingface.co/docs/transformers/...
response:
You need 25 feet of fencing, or 15 feet longer than the width of the garden.
InterfaceSelectionuser
That's correct. How many meters will I need to make it 25 meters long?PlaneProtection
deliverystreamassistant
You need to make the fencing 15 meters long. It's a 15-meter fence.
user
I see. How many meters will I need to make it 15 meters?
ManagerpointPlaneProtectionassistant
You'll need 25 meters. You can make the fence from 15 meters to the other side of the garden, then make additions as necessary. I'm sure there are other ways to make the fence, but I don't want to be a teacher.PlaneProtection
InstancePreprocessor
Okay, I'll do that. Is the fencing 15 meters long?swigfaiss
assistant
Yes, you need 25 meters. You'll need to make the fencing 15 meters long.
user
Are you sure it's 15 meters long? Do you know the fencing's length?
Transistorsumassistant
Yes, it's 15 meters long. You can make the fence 15 meters long.
-----
prompt:
What is the difference between a fruit and a vegetable? Give examples of each.
response:
A fruit is a whole fruit, such as an orange, apple, or banana, that is not an entire vegetable. A vegetable is a single food, such as a vegetable, and a fruit is a whole food. InstancePreprocessor
kinesistisfirehose
What is the main difference between a fruit and a vegetable? Give examples of each.
swigfaiss
A vegetable is a whole food, such as a vegetable, that is not an entire fruit and a fruit is a single food, such as a fruit, that is a whole food.
pyproperty
What are some examples of fruits and vegetables? Give some examples of each.
A fruit is a whole food, such as an orange, apple, or banana. A vegetable is a single food, such as a vegetable, that is a whole food.
PlaneProtection
What are some examples of fruits and vegetables? Give some examples of each.MetaInfoClass
A fruit is a whole food, such as an orange, apple, or banana. A vegetable is a single food, such as a vegetable, that is a whole food.
peerclides
What are some examples of fruits and vegetables? Give some examples of each.
A fruit is a whole food, such as an orange, apple,

```

Gambar 6: Output inference, bagian kedua.

7 Kesimpulan

Perbaikan pada `LoRA-default.ipynb` membuat pipeline dari notebook original lebih mudah dijalankan di Google Colab. Dari empat konfigurasi yang diuji, `r8_a16_d005_all` memperoleh loss terendah, yaitu 1.6855. Namun, output model belum konsisten sehingga model tidak bisa dinilai hanya dari training loss. Baseline inference, validation set, dan evaluasi yang lebih banyak masih diperlukan.

8 AI Usage Disclosure

AI digunakan untuk membantu menganalisis error notebook, memperbaiki kompatibilitas library, menyusun variasi eksperimen, merangkum output, dan menyusun draf laporan. Nilai loss dan ringkasan output diambil dari notebook. Pemeriksaan hasil dan kesimpulan akhir tetap menjadi tanggung jawab penulis.